

Playing with a small DB and storage

Setup visualization libraries

```
In [2]: def displaySectionCaption(caption, color='coral'):
        html_string = f'<hr><strong><p style="color:{color};font-size:16px;">{caption}</p.</strong>'
        display(HTML(html_string))
```

We study a simplified IO model for HDDs and SSDs in CMPE-138. The model will work well in practice, for our query optimization and data layout problems.

```
In [4]: import math
        from math import ceil, log

        # We'll use MBs -- for basic i to MBs
        (MB, GB, TB, KB, Bytes) = (1.0, 1024.0, 1024.0*1024.0,
                                     1.0/1024.0, 1.0/(1024.0*1024))

        # 64 MB-Blocks (default)
        PageSizeMB = 64.0*MB
        size_of_types = {'int64': 8, 'int32': 4, 'double': 8, 'char': 1} # in bytes

        class IOdevice:
            def __init__(self, accessTime, scanSpeed, C_w):
                self.C_r = 1.0 # Cost of reads
                self.C_w = C_w # Cost of writes relative to reads
                self.accessTime = accessTime
                self.scanSpeed = scanSpeed

            # Read costs: Simple IOcost model using Access time + Scan speeds
            def read_pages_cost(self, numPages):
                # Assume you need to read full pages. (i.e., no partial pages)
                numPages = math.ceil(numPages)
                tsecs = numPages*self.accessTime # time to access
                tsecs += numPages*PageSizeMB/self.scanSpeed # time to scan
                return (tsecs)

            def write_pages_cost(self, numPages):
                return self.C_w*self.read_pages_cost(numPages)

        # Example IO devices in 2024
        # Access and Scan speeds in [seconds, MBps], Cw cost of write vs reads.
        ram1 = IOdevice(100*pow(10, -9), 100.0*1024, 1.0)
        ssd1 = IOdevice(10*pow(10, -6), 5.0*1024, 1.0) # 10 microsecs, 5GBps
        hdd1 = IOdevice(10*pow(10, -3), 100.0, 1.0) # 10 millisecs, 100 MBps
        # machine to machine over network (modeling a network as an IO device)
        m2m1 = IOdevice(10*pow(10, -6), 5.0*1024, 1.0) # 1 micro, 5 GBps

        IOdevices1 = {'HDD': hdd1, 'SSD': ssd1, 'RAM': ram1}
```

```
In [6]: """
        Basic physical table
        """
        class Table:
            def __init__(self, sizeInMBs, rowSize):
                self.sizeInMBs = sizeInMBs
                self.rowSize = rowSize
                self.numRows = ceil(self.sizeInMBs/self.rowSize)

            # self.numTuples = numTuples
            self.isSorted = False
            self.isHPed = False

            # P(R) -- number of Pages for table
            def P(self):
                P = ceil(self.sizeInMBs/PageSizeMB)
                return P
            def RowSize(self):
                return self.rowSize
            def T(self):
                return self.numRows
            def SizeInMBs(self):
                return self.sizeInMBs

            # Keeping track of is table sorted, HPed, or neither (default)
            def Sort(self):
                self.isSorted = True
                self.isHPed = False
            def HP(self):
                self.isSorted = False
                self.isHPed = True
            def Reset(self):
                self.isSorted = False
                self.isHPed = False
```

Exercises:

```
In [12]: # Spotify Songs Table [songid: int64, title: text, name: text, genre: text]
        # -- Size of row = 8 bytes (int64) + avg size of title+name+genre.
        # -- Assume avg row size = 1024 Bytes
        songs_rowSize = 1024.0*Bytes
        songs_numRows = 500000000.0 # 500 million songs

        """Problem 1:
        Calculate the size (MBs) of SongsTable, and num pages."""
        songsTableSize = songs_rowSize * songs_numRows / (1024 * 1024) # Convert to MBs
        songsNumPages = math.ceil(songsTableSize / PageSizeMB)
        print(f"Songs Table Size: {songsTableSize:.2f} MB, Number of Pages: {songsNumPages}")

        """Problem 2: Read costs
        Compute the cost in seconds to read 100 pages from the SongsTable"""
        numPagesToRead = 100
        hdd_read_cost = hdd1.read_pages_cost(numPagesToRead)
        ssd_read_cost = ssd1.read_pages_cost(numPagesToRead)
        ram_read_cost = ram1.read_pages_cost(numPagesToRead)
        print(f"Read cost (HDD): {hdd_read_cost:.6f} sec")
        print(f"Read cost (SSD): {ssd_read_cost:.6f} sec")
        print(f"Read cost (RAM): {ram_read_cost:.6f} sec")

        """Problem 3: Effect of caching
        Read 200 pages. 1st check RAM.
        - Cache hit of 90% in RAM.
        - For RAM cache misses (the other 10%), 75% are in SSD and 25% are in HDD."""
        numPagesToRead = 200
        cache_hit_ratio = 0.90
        ssd_miss_ratio = 0.75
        hdd_miss_ratio = 0.25

        ram_hits = cache_hit_ratio * numPagesToRead
        ram_misses = numPagesToRead - ram_hits
        ssd_hits = ssd_miss_ratio * ram_misses
        hdd_hits = hdd_miss_ratio * ram_misses

        total_read_time = (
            ram1.read_pages_cost(ram_hits) +
            ssd1.read_pages_cost(ssd_hits) +
            hdd1.read_pages_cost(hdd_hits)
        )

        print(f"Total read time with caching: {total_read_time:.6f} sec")
```

Songs Table Size: 0.47 MB, Number of Pages: 1
Read cost (HDD): 65.000000 sec
Read cost (SSD): 1.251000 sec
Read cost (RAM): 0.062510 sec
Total read time with caching: 3.550168 sec

In []: