

# CMPE 138/180B Database Systems

Midterm review  
Dr. Gheorghi Guzun

# Midterm Exam: 3/27

---

- During class time
- Closed book. One letter sized page of hand written notes allowed
- Calculator allowed
- Duration: 1hr 10min
- Combination of multiple choice, short answer, SQL queries, and problem solving exercises.

# Typical DBMS Functionality

- *Define* a particular database in terms of its data types, structures, and constraints
- *Construct* or Load the initial database contents on a secondary storage medium
- *Manipulating* the database:
  - Retrieval: Querying, generating reports
  - Modification: Insertions, deletions and updates to its content
  - Accessing the database
- *Processing* and *Sharing* by a set of concurrent users and application programs – yet, keeping all data valid and consistent

# Application Activities Against a Database

- Applications interact with a database by generating
  - Queries: that access different parts of data and formulate the result of a request
  - Transactions: that may read some data and “update” certain values or generate new data and store that in the database
- Applications must not allow unauthorized users to access data
- Applications must keep up with changing user requirements against the database

# Example of a simple database

## COURSE

Course_name	Course_number	Credit_hours	Department
Intro to Computer Science	CS1310	4	CS
Data Structures	CS3320	4	CS
Discrete Mathematics	MATH2410	3	MATH
Database	CS3380	3	CS

## SECTION

Section_identifier	Course_number	Semester	Year	Instructor
85	MATH2410	Fall	04	King
92	CS1310	Fall	04	Anderson
102	CS3320	Spring	05	Knuth
112	MATH2410	Fall	05	Chang
119	CS1310	Fall	05	Anderson
135	CS3380	Fall	05	Stone

## GRADE\_REPORT

Student_number	Section_identifier	Grade
17	112	B
17	119	C
8	85	A
8	92	A
8	102	B
8	135	A

## PREREQUISITE

Course_number	Prerequisite_number
CS3380	CS3320
CS3380	MATH2410
CS3320	CS1310

**Figure 1.2**

A database that stores student and course information.

# Main Characteristics of the Database Approach

- **Self-describing nature of a database system:**

- A DBMS **catalog** stores the description of a particular database (e.g. data structures, types, and constraints)
- The description is called **meta-data**\*
- This allows the DBMS software to work with different database applications.

- **Insulation between programs and data:**

- Called **program-data independence**.
- Allows changing data structures and storage organization without having to change the DBMS access programs.

-----

\* Some newer systems such as a few NOSQL systems need no meta-data: they store the data definition within its structure making it self describing

# Database Schema vs. Database State (continued)

- Distinction
  - The ***database schema*** changes very infrequently.
  - The ***database state*** changes every time the database is updated.
- Schema is also called **intension**.
- State is also called **extension**.

# Example of a Database Schema

## STUDENT

Name	Student_number	Class	Major
------	----------------	-------	-------

## COURSE

Course_name	Course_number	Credit_hours	Department
-------------	---------------	--------------	------------

## PREREQUISITE

Course_number	Prerequisite_number
---------------	---------------------

## SECTION

Section_identifier	Course_number	Semester	Year	Instructor
--------------------	---------------	----------	------	------------

## GRADE\_REPORT

Student_number	Section_identifier	Grade
----------------	--------------------	-------

**Figure 2.1**

Schema diagram for the database in Figure 1.2.



# Example of a database state

## COURSE

Course_name	Course_number	Credit_hours	Department
Intro to Computer Science	CS1310	4	CS
Data Structures	CS3320	4	CS
Discrete Mathematics	MATH2410	3	MATH
Database	CS3380	3	CS

## SECTION

Section_identifier	Course_number	Semester	Year	Instructor
85	MATH2410	Fall	04	King
92	CS1310	Fall	04	Anderson
102	CS3320	Spring	05	Knuth
112	MATH2410	Fall	05	Chang
119	CS1310	Fall	05	Anderson
135	CS3380	Fall	05	Stone

## GRADE\_REPORT

Student_number	Section_identifier	Grade
17	112	B
17	119	C
8	85	A
8	92	A
8	102	B
8	135	A

## PREREQUISITE

Course_number	Prerequisite_number
CS3380	CS3320
CS3380	MATH2410
CS3320	CS1310

**Figure 1.2**

A database that stores student and course information.

# Example of a Relation

STUDENT						
Name	Ssn	Home_phone	Address	Office_phone	Age	Gpa
Benjamin Bayer	305-61-2435	373-1616	2918 Bluebonnet Lane	NULL	19	3.21
Chung-cha Kim	381-62-1245	375-4409	125 Kirby Road	NULL	18	2.89
Dick Davidson	422-11-2320	NULL	3452 Elgin Road	749-1253	25	3.53
Rohan Panchal	489-22-1100	376-9821	265 Lark Lane	749-6492	28	3.93
Barbara Benson	533-69-1238	839-8461	7384 Fontana Lane	NULL	19	3.25

**Figure 5.1**

The attributes and tuples of a relation STUDENT.

# Formal Definitions - Schema

- The **Schema** (or description) of a Relation:
  - Denoted by  $R(A_1, A_2, \dots, A_n)$
  - $R$  is the **name** of the relation
  - The **attributes** of the relation are  $A_1, A_2, \dots, A_n$
- Example:  
CUSTOMER (Cust-id, Cust-name, Address, Phone#)
  - CUSTOMER is the relation name
  - Defined over the four attributes: Cust-id, Cust-name, Address, Phone#
- Each attribute has a **domain** or a set of valid values.
  - For example, the domain of Cust-id is 6 digit numbers.

# Formal Definitions - Tuple

- A **tuple** is an ordered set of values (enclosed in angled brackets '< ... >')
- Each value is derived from an appropriate *domain*.
- A row in the CUSTOMER relation is a 4-tuple and would consist of four values, for example:
  - <632895, "John Smith", "101 Main St. Atlanta, GA 30332", "(404) 894-2000">
  - This is called a 4-tuple as it has 4 values
  - A tuple (row) in the CUSTOMER relation.
- A relation is a **set** of such tuples (rows)

# Formal Definitions - Domain

- A **domain** has a logical definition:
  - Example: “USA\_phone\_numbers” are the set of 10 digit phone numbers valid in the U.S.
- A domain also has a data-type or a format defined for it.
  - The USA\_phone\_numbers may have a format: (ddd)ddd-dddd where each d is a decimal digit.
  - Dates have various formats such as year, month, date formatted as yyyy-mm-dd, or as dd mm,yyyy etc.
- The attribute name designates the role played by a domain in a relation:
  - Used to interpret the meaning of the data elements corresponding to that attribute
  - Example: The domain Date may be used to define two attributes named “Invoice-date” and “Payment-date” with different meanings

# Formal Definitions - State

- The **relation state** is a subset of the Cartesian product of the domains of its attributes
  - each domain contains the set of all possible values the attribute can take.
- Example: attribute Cust-name is defined over the domain of character strings of maximum length 25
  - $\text{dom}(\text{Cust-name})$  is `varchar(25)`
- The role these strings play in the CUSTOMER relation is that of the *name of a customer*.

# Formal Definitions - Summary

- Formally,
  - Given  $R(A_1, A_2, \dots, A_n)$
  - $r(R) \subset \text{dom}(A_1) \times \text{dom}(A_2) \times \dots \times \text{dom}(A_n)$
- $R(A_1, A_2, \dots, A_n)$  is the **schema** of the relation
- $R$  is the **name** of the relation
- $A_1, A_2, \dots, A_n$  are the **attributes** of the relation
- $r(R)$ : a specific **state** (or "value" or "population") of relation  $R$  – this is a *set of tuples* (rows)
  - $r(R) = \{t_1, t_2, \dots, t_n\}$  where each  $t_i$  is an  $n$ -tuple
  - $t_i = \langle v_1, v_2, \dots, v_n \rangle$  where each  $v_j$  *element-of*  $\text{dom}(A_j)$

# Definition Summary

<u>Informal Terms</u>		<u>Formal Terms</u>
Table		Relation
Column Header		Attribute
All possible Column Values		Domain
Row		Tuple
Table Definition		Schema of a Relation
Populated Table		State of the Relation



# Relational Integrity Constraints

- Constraints are **conditions** that must hold on **all** valid relation states.
- There are three *main types* of (explicit schema-based) constraints that can be expressed in the relational model:
  - **Key** constraints
  - **Entity integrity** constraints
  - **Referential integrity** constraints
- Another schema-based constraint is the **domain** constraint
  - Every value in a tuple must be from the *domain of its attribute* (or it could be **null**, if allowed for that attribute)

# COMPANY Database Schema

## EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
-------	-------	-------	------------	-------	---------	-----	--------	-----------	-----

## DEPARTMENT

Dname	<u>Dnumber</u>	Mgr_ssn	Mgr_start_date
-------	----------------	---------	----------------

## DEPT\_LOCATIONS

<u>Dnumber</u>	<u>Dlocation</u>
----------------	------------------

## PROJECT

Pname	<u>Pnumber</u>	Plocation	Dnum
-------	----------------	-----------	------

## WORKS\_ON

<u>Essn</u>	<u>Pno</u>	Hours
-------------	------------	-------

## DEPENDENT

<u>Essn</u>	<u>Dependent_name</u>	Sex	Bdate	Relationship
-------------	-----------------------	-----	-------	--------------

**Figure 5.5**  
Schema diagram for  
the COMPANY  
relational database  
schema.

# Populated database state for COMPANY

**Figure 5.6**

One possible database state for the COMPANY relational database schema.

## EMPLOYEE

Fname	Minit	Lname	Ssn	Bdate	Address	Sex	Salary	Super_ssn	Dno
John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
Alicia	J	Zelaya	999887777	1968-01-19	3321 Castle, Spring, TX	F	25000	987654321	4
Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	NULL	1

## DEPARTMENT

Dname	Dnumber	Mgr_ssn	Mgr_start_date
Research	5	333445555	1988-05-22
Administration	4	987654321	1995-01-01
Headquarters	1	888665555	1981-06-19

## DEPT\_LOCATIONS

Dnumber	Dlocation
1	Houston
4	Stafford
5	Bellaire
5	Sugarland
5	Houston

## WORKS\_ON

Essn	Pno	Hours
123456789	1	32.5
123456789	2	7.5
666884444	3	40.0
453453453	1	20.0
453453453	2	20.0
333445555	2	10.0
333445555	3	10.0
333445555	10	10.0
333445555	20	10.0
999887777	30	30.0
999887777	10	10.0
987987987	10	35.0
987987987	30	5.0
987654321	30	20.0
987654321	20	15.0
888665555	20	NULL

## PROJECT

Pname	Pnumber	Plocation	Dnum
ProductX	1	Bellaire	5
ProductY	2	Sugarland	5
ProductZ	3	Houston	5
Computerization	10	Stafford	4
Reorganization	20	Houston	1
Newbenefits	30	Stafford	4

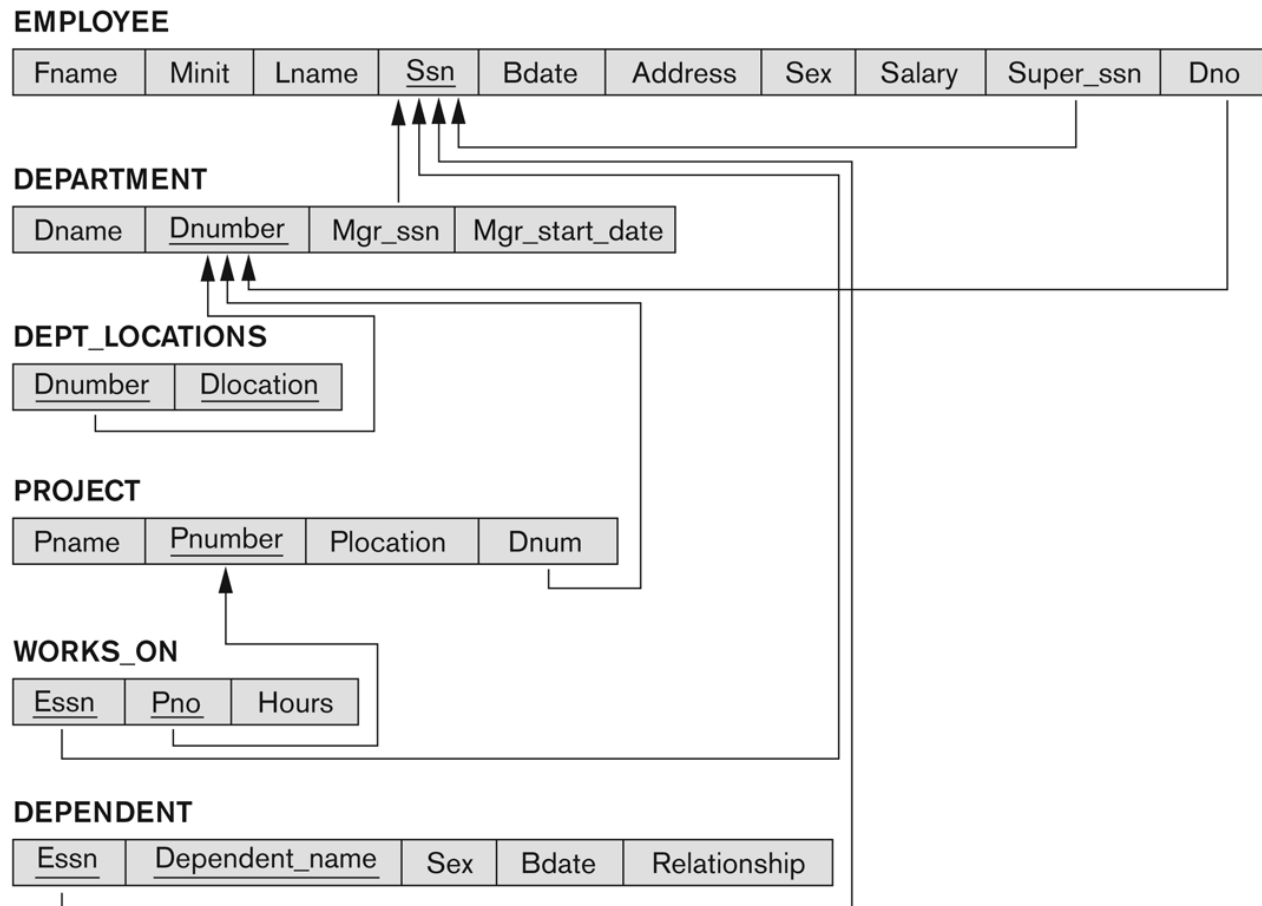
## DEPENDENT

Essn	Dependent_name	Sex	Bdate	Relationship
333445555	Alice	F	1986-04-05	Daughter
333445555	Theodore	M	1983-10-25	Son
333445555	Joy	F	1958-05-03	Spouse
987654321	Abner	M	1942-02-28	Spouse
123456789	Michael	M	1988-01-04	Son
123456789	Alice	F	1988-12-30	Daughter
123456789	Elizabeth	F	1967-05-05	Spouse

# Referential Integrity Constraints for COMPANY database

### Figure 5.7

Referential integrity constraints displayed on the COMPANY relational database schema.



# Basic SQL Queries,

- SQL Data Definition and Data Types
- Specifying Constraints in SQL
- Basic Retrieval Queries in SQL
- INSERT, DELETE, and UPDATE Statements in SQL
- Additional Features of SQL

# The CREATE TABLE Command in SQL

- Specifying a new relation
  - Provide name of table
  - Specify attributes, their types and initial constraints
- Can optionally specify schema:
  - `CREATE TABLE COMPANY.EMPLOYEE ...`  
or
  - `CREATE TABLE EMPLOYEE ...`

# Attribute Data Types and Domains in SQL

- **Basic data types**

- **Numeric data types**

- Integer numbers: INTEGER, INT, and SMALLINT
    - Floating-point (real) numbers: FLOAT or REAL, and DOUBLE PRECISION

- **Character-string data types**

- Fixed length: CHAR(*n*), CHARACTER(*n*)
    - Varying length: VARCHAR(*n*), CHAR VARYING(*n*), CHARACTER VARYING(*n*)

# Attribute Data Types and Domains in SQL (cont'd.)

- **Bit-string** data types
  - Fixed length: `BIT (n)`
  - Varying length: `BIT VARYING (n)`
- **Boolean** data type
  - Values of `TRUE` or `FALSE` or `NULL`
- **DATE** data type
  - Ten positions
  - Components are `YEAR`, `MONTH`, and `DAY` in the form `YYYY-MM-DD`
  - Multiple mapping functions available in RDBMSs to change date formats



# Attribute Data Types and Domains in SQL (cont'd.)

- Additional data types

- **Timestamp** data type

Includes the `DATE` and `TIME` fields

- Plus a minimum of six positions for decimal fractions of seconds
    - Optional `WITH TIME ZONE` qualifier

- **INTERVAL** data type

- Specifies a relative value that can be used to increment or decrement an absolute value of a date, time, or timestamp

- **DATE, TIME, Timestamp, INTERVAL** data types can be **cast** or converted to string formats for comparison.

# Attribute Data Types and Domains in SQL (cont'd.)

## ■ Domain

- Name used with the attribute specification
- Makes it easier to change the data type for a domain that is used by numerous attributes
- Improves schema readability
- Example:
  - `CREATE DOMAIN SSN_TYPE AS CHAR(9);`

## ■ TYPE

- User Defined Types (UDTs) are supported for object-oriented applications. (See Ch.12) Uses the command: `CREATE TYPE`

# Specifying Constraints in SQL

## Basic constraints:

- Relational Model has 3 basic constraint types that are supported in SQL:
  - **Key** constraint: A primary key value cannot be duplicated
  - **Entity Integrity** Constraint: A primary key value cannot be null
  - **Referential integrity** constraints : The “foreign key” must have a value that is already present as a primary key, or may be null.

# Specifying Attribute Constraints

Other Restrictions on attribute domains:

- Default value of an attribute
  - **DEFAULT** <value>
  - **NULL** is not permitted for a particular attribute (**NOT NULL**)
- **CHECK** clause
  - `Dnumber INT NOT NULL CHECK (Dnumber > 0 AND Dnumber < 21);`

# Specifying Key and Referential Integrity Constraints

## ■ PRIMARY KEY clause

- Specifies one or more attributes that make up the primary key of a relation
- `Dnumber INT PRIMARY KEY;`

## ■ UNIQUE clause

- Specifies alternate (secondary) keys (called CANDIDATE keys in the relational model).
- `Dname VARCHAR(15) UNIQUE;`

# Specifying Key and Referential Integrity Constraints (cont'd.)

- **FOREIGN KEY** clause
  - Default operation: reject update on violation
  - Attach **referential triggered action** clause
    - Options include SET NULL, CASCADE, and SET DEFAULT
    - Action taken by the DBMS for SET NULL or SET DEFAULT is the same for both ON DELETE and ON UPDATE
    - CASCADE option suitable for “relationship” relations

# Basic SQL Retrieval Query Block

```
SELECT    <attribute list>  
FROM      <table list>  
[ WHERE    <condition> ]  
[ ORDER BY <attribute list> ];
```

# INSERT, DELETE, and UPDATE Statements in SQL

- Three commands used to modify the database:
  - INSERT, DELETE, and UPDATE
- INSERT typically inserts a tuple (row) in a relation (table)
- UPDATE may update a number of tuples (rows) in a relation (table) that satisfy the condition
- DELETE may also update a number of tuples (rows) in a relation (table) that satisfy the condition



# Nested Queries (cont'd.)

```
Q4A:  SELECT  DISTINCT Pnumber
      FROM    PROJECT
      WHERE   Pnumber IN
            ( SELECT  Pnumber
              FROM    PROJECT, DEPARTMENT, EMPLOYEE
              WHERE   Dnum=Dnumber AND
                    Mgr_ssn=Ssn AND Lname='Smith' )

      OR

      Pnumber IN
            ( SELECT  Pno
              FROM    WORKS_ON, EMPLOYEE
              WHERE   Essn=Ssn AND Lname='Smith' );
```

# USE of EXISTS

**Q7:**

```
SELECT Fname, Lname
FROM Employee
WHERE EXISTS (SELECT *
               FROM DEPENDENT
               WHERE Ssn= Essn)

      AND EXISTS (SELECT *
                  FROM Department
                  WHERE Ssn= Mgr_Ssn)
```

# USE of EXISTS

**Q7:**

```
SELECT Fname, Lname
FROM Employee
WHERE EXISTS (SELECT *
               FROM DEPENDENT
               WHERE Ssn= Essn)

      AND EXISTS (SELECT *
                  FROM Department
                  WHERE Ssn= Mgr_Ssn)
```

# How to think about SELECT versus logic

1. Output versus logic
  - a. SELECT specifies the output (like “return(a, b)” in a function)
  - b. GROUP BY, ON, ORDER ... control the logic
2. After you GROUP BY a set of columns (e.g., title, artist), the extra columns (e.g., genre) are not available anymore for the subsequent logic.

The diagram shows a SQL query with three components highlighted and labeled:

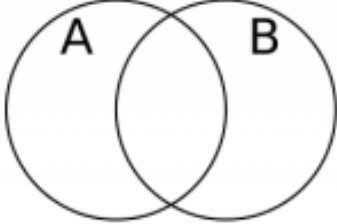
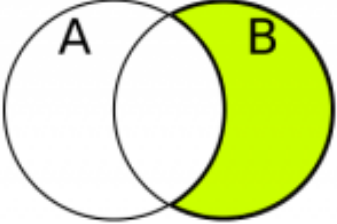
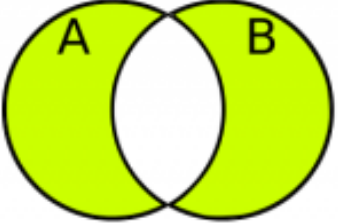
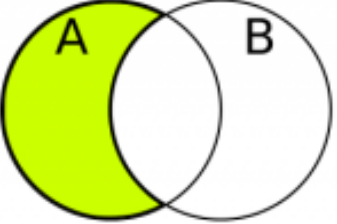
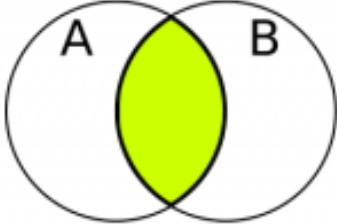
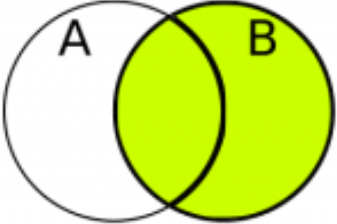
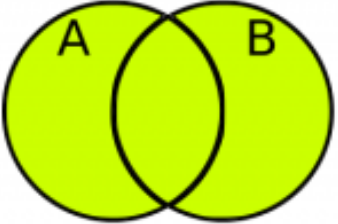
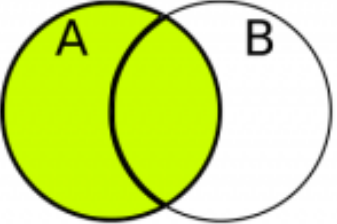
- OUTPUT:** Points to the `SELECT` clause, which is enclosed in a red box. The text `SELECT songs.title, songs.artist, count(Listens.song_id) as popular` is highlighted in blue.
- INPUT:** Points to the `FROM` and `JOIN` clauses, which are enclosed in a green box. The text `FROM Songs JOIN Listens` is highlighted in green.
- LOGIC:** Points to the `ON`, `GROUP BY`, `ORDER BY`, and `LIMIT` clauses, which are enclosed in a purple box. The text `ON Songs.song_id = Listens.song_id GROUP BY songs.title, songs.artist ORDER BY COUNT(Listens.song_id) DESC LIMIT 10;` is highlighted in purple.

```
-- Popular songs based on number of listens
SELECT songs.title, songs.artist, count(Listens.song_id) as popular
FROM Songs
JOIN Listens
ON Songs.song_id = Listens.song_id
GROUP BY songs.title, songs.artist
ORDER BY COUNT(Listens.song_id) DESC
LIMIT 10;
```

# Different Types of JOINed Tables in SQL

- Specify different types of join
  - NATURAL JOIN
  - Various types of OUTER JOIN (LEFT, RIGHT, FULL )
- NATURAL JOIN on two relations R and S
  - No join condition specified
  - Is equivalent to an implicit EQUIJOIN condition for each pair of attributes with same name from R and S

# Selecting from multiple tables: SQL JOINS

SQL JOINS			
Arranged in a Karnaugh Map by Jason Charney (jrcharney@gmail.com)			
No join ( $\emptyset$ )	[Exclusive] Right Join ( $\sim A$ )	[Exclusive] Full Join ( $A \oplus B$ )	[Exclusive] Left Join ( $\sim B$ )
	 <pre>SELECT * FROM A RIGHT JOIN B ON A.key = B.key WHERE B.key IS NULL</pre>	 <pre>SELECT * FROM A FULL JOIN B ON A.key = B.key WHERE B.key IS NULL OR A.key IS NULL</pre>	 <pre>SELECT * FROM A LEFT JOIN B ON A.key = B.key WHERE B.key IS NULL</pre>
Inner Join ( $A \wedge B$ )	[Inclusive] Right Join ( $B$ )	[Inclusive] Full Join ( $A \vee B$ )	[Inclusive] Left Join ( $A$ )
 <pre>SELECT * FROM A INNER JOIN B ON A.key = B.key</pre>	 <pre>SELECT * FROM A RIGHT JOIN B ON A.key = B.key</pre>	 <pre>SELECT * FROM A FULL JOIN B ON A.key = B.key</pre>	 <pre>SELECT * FROM A LEFT JOIN B ON A.key = B.key</pre>

# Aggregate Functions in SQL

- Used to summarize information from multiple tuples into a single-tuple summary
- Built-in aggregate functions
  - **COUNT**, **SUM**, **MAX**, **MIN**, and **AVG**
- **Grouping**
  - Create subgroups of tuples before summarizing
- To select entire groups, **HAVING** clause is used
- Aggregate functions can be used in the **SELECT** clause or in a **HAVING** clause

# Grouping: The GROUP BY Clause

- **Partition** relation into subsets of tuples
  - Based on **grouping attribute(s)**
  - Apply function to each such group independently
- **GROUP BY** clause
  - Specifies grouping attributes
- **COUNT (\*)** counts the number of rows in the group



# EXPANDED Block Structure of SQL Queries

```
SELECT <attribute and function list>  
FROM <table list>  
[ WHERE <condition> ]  
[ GROUP BY <grouping attribute(s)> ]  
[ HAVING <group condition> ]  
[ ORDER BY <attribute list> ];
```

# Table 7.2 Summary of SQL Syntax

**Table 7.2** Summary of SQL Syntax

---

```
CREATE TABLE <table name> ( <column name> <column type> [ <attribute constraint> ]  
                             { , <column name> <column type> [ <attribute constraint> ] }  
                             [ <table constraint> { , <table constraint> } ] )
```

---

```
DROP TABLE <table name>  
ALTER TABLE <table name> ADD <column name> <column type>
```

---

```
SELECT [ DISTINCT ] <attribute list>  
FROM ( <table name> { <alias> } | <joined table> ) { , ( <table name> { <alias> } | <joined table> ) }  
[ WHERE <condition> ]  
[ GROUP BY <grouping attributes> [ HAVING <group selection condition> ] ]  
[ ORDER BY <column name> [ <order> ] { , <column name> [ <order> ] } ]
```

---

```
<attribute list> ::= ( * | ( <column name> | <function> ( ( [ DISTINCT ] <column name> | * ) ) )  
                    { , ( <column name> | <function> ( ( [ DISTINCT ] <column name> | * ) ) ) } ) )
```

---

```
<grouping attributes> ::= <column name> { , <column name> }
```

---

```
<order> ::= ( ASC | DESC )
```

---

```
INSERT INTO <table name> [ ( <column name> { , <column name> } ) ]  
( VALUES ( <constant value> , { <constant value> } ) { , ( <constant value> { , <constant value> } ) }  
| <select statement> )
```

---

*continued on next slide*

# Table 7.2 (continued)

## Summary of SQL Syntax

**Table 7.2** Summary of SQL Syntax

---

DELETE FROM <table name>

[ WHERE <selection condition> ]

---

UPDATE <table name>

SET <column name> = <value expression> { , <column name> = <value expression> }

[ WHERE <selection condition> ]

---

CREATE [ UNIQUE] INDEX <index name>

ON <table name> ( <column name> [ <order> ] { , <column name> [ <order> ] } )

[ CLUSTER ]

---

DROP INDEX <index name>

---

CREATE VIEW <view name> [ ( <column name> { , <column name> } ) ]

AS <select statement>

---

DROP VIEW <view name>

---

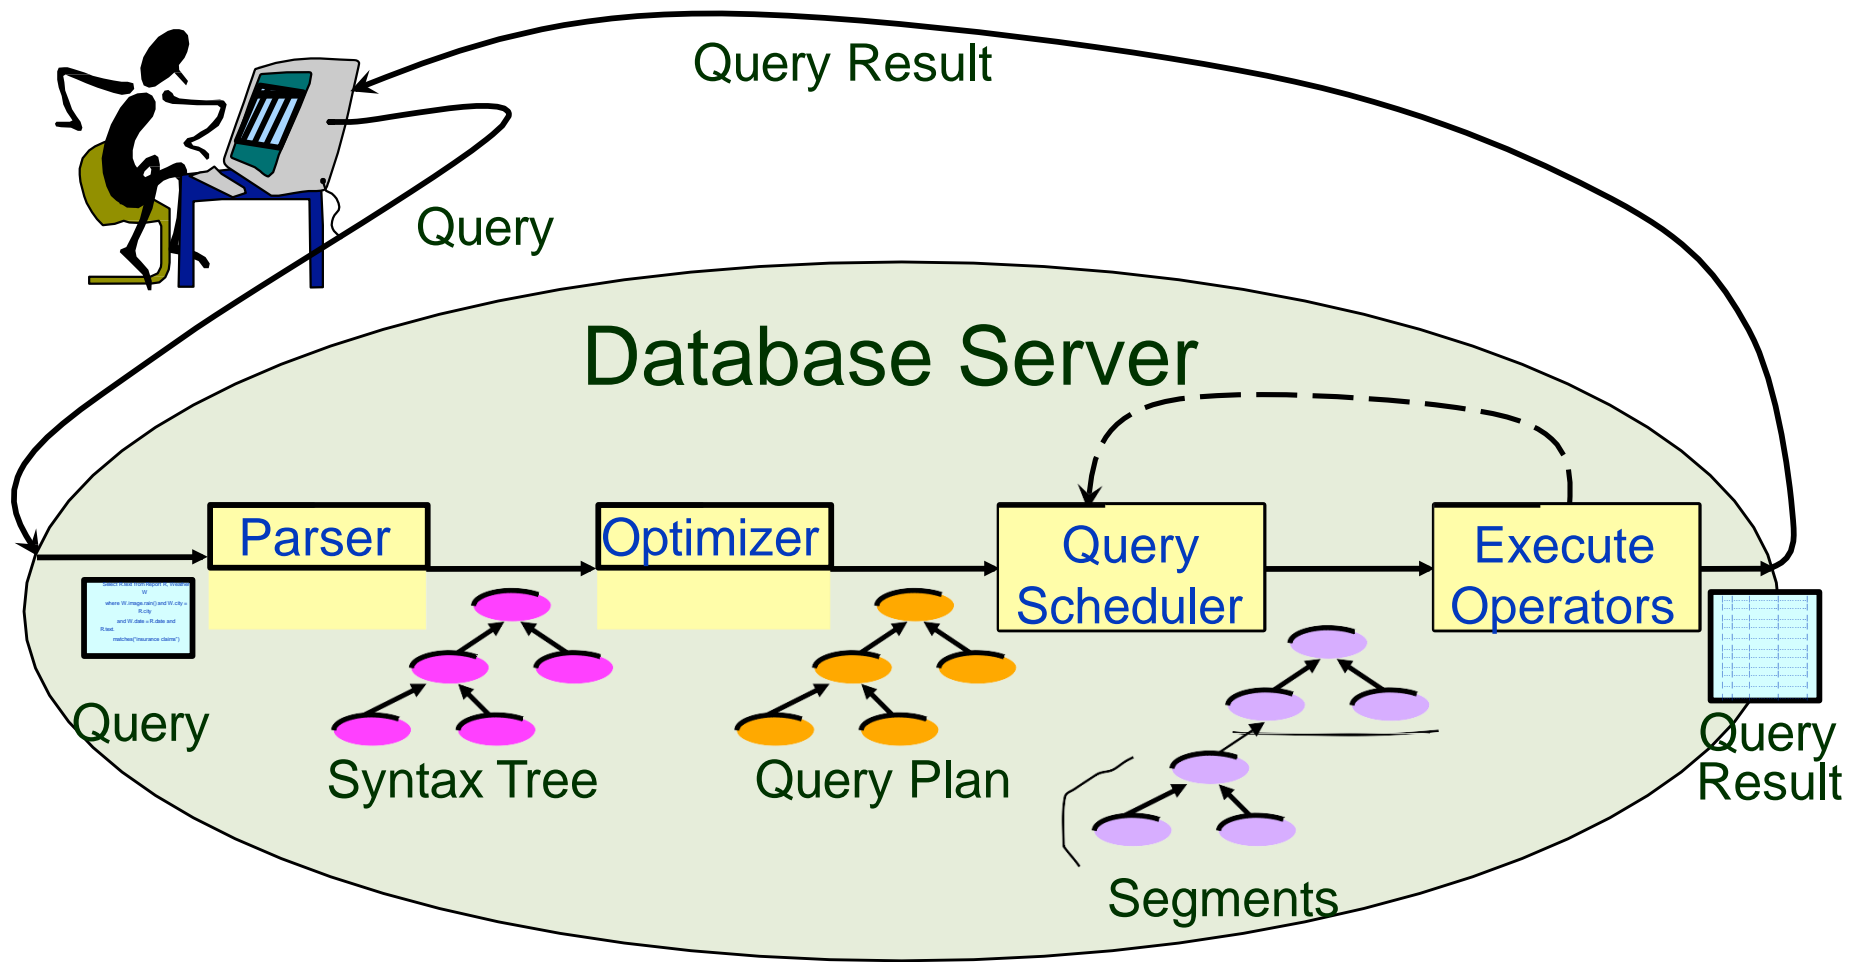
NOTE: The commands for creating and dropping indexes are not part of standard SQL.

# Database Management System (DBMS)


- A DBMS *manages* a *database*.
- A database is a collection of data, usually with some description of the structure of the data.
  - Structure description, if present, is described using a *schema*. e.g. the CREATE TABLE command in SQL



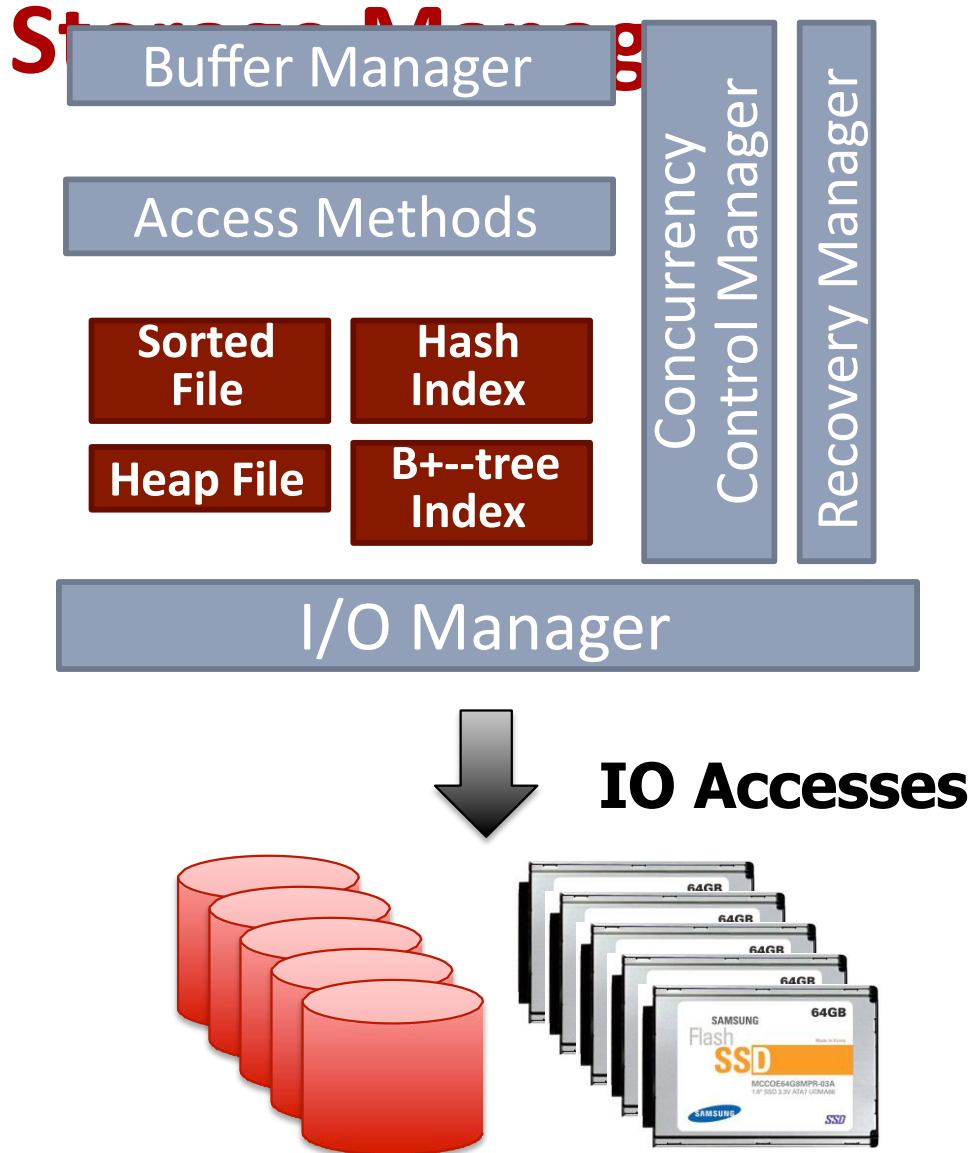
# Life Cycle of a Query



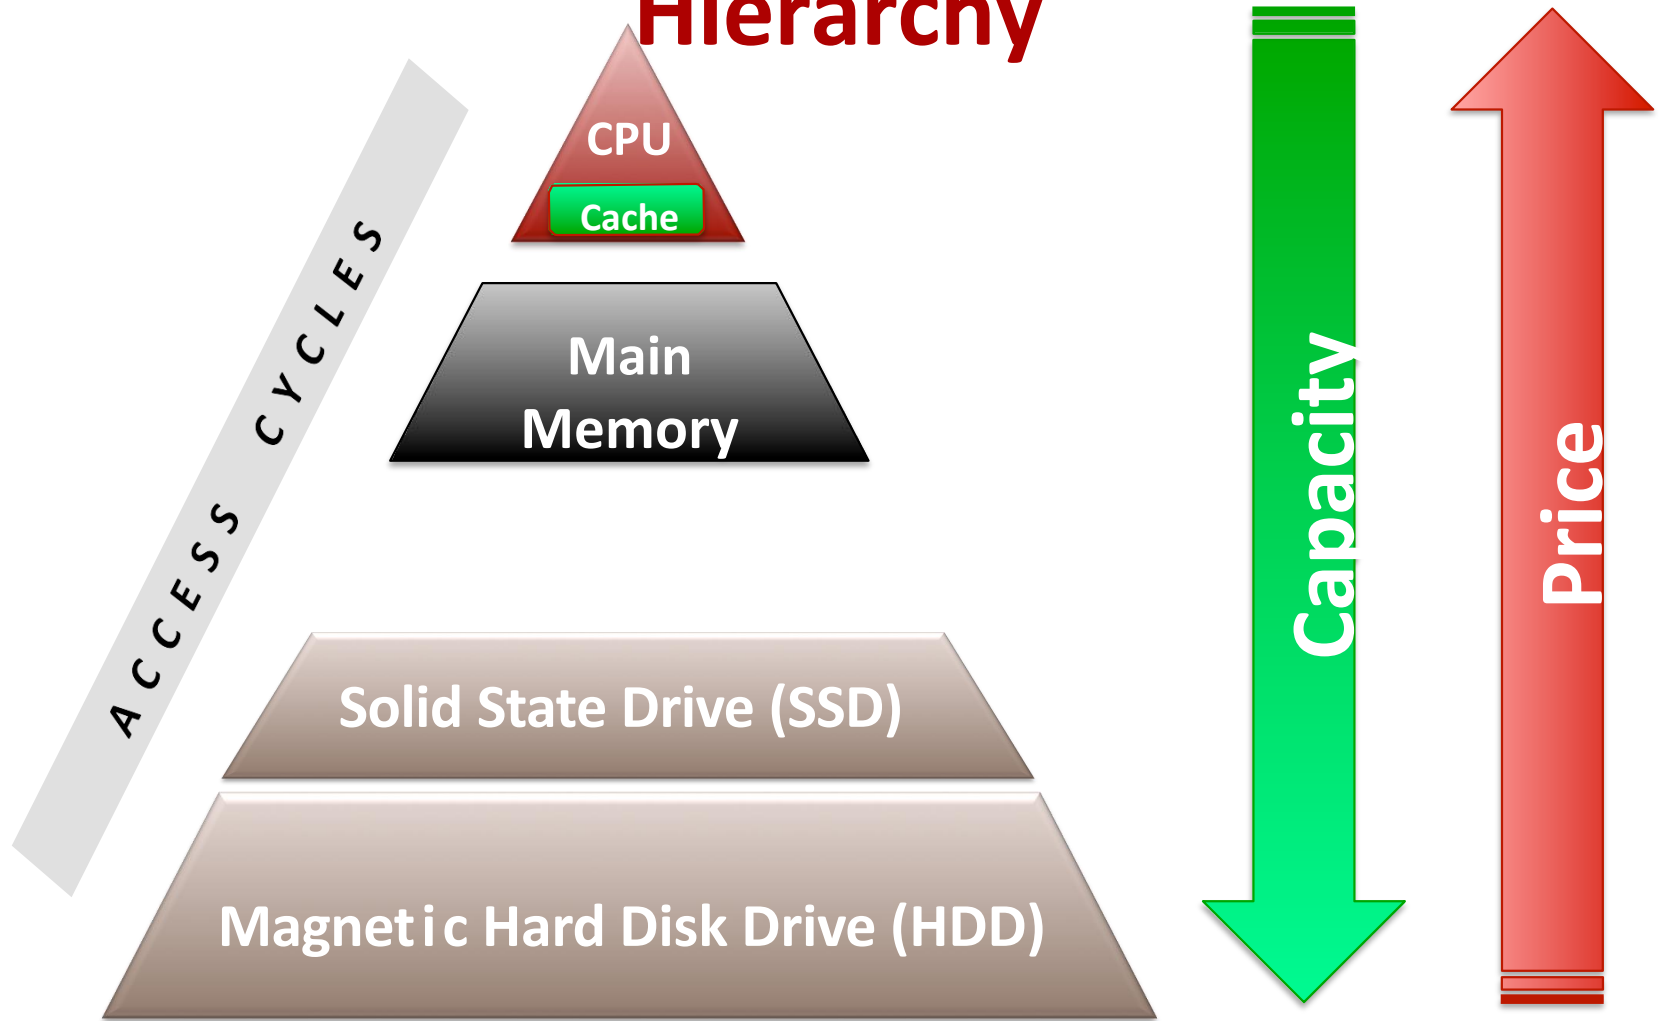
# Internal Architecture of a Data Processing System



# Architecture of a



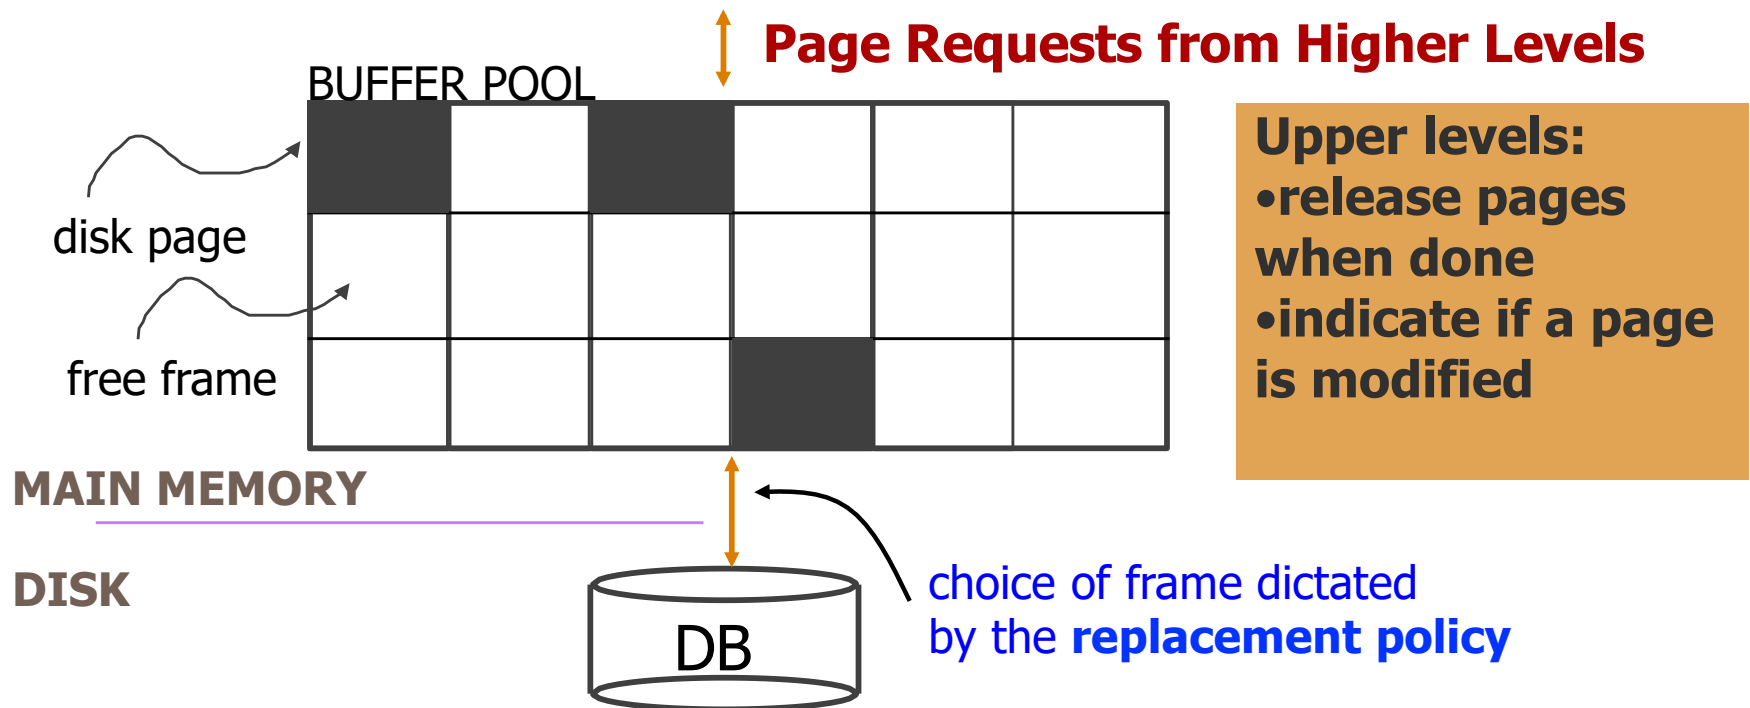
# Memory Hierarchy





# Buffer Management in a DBMS

- Data must be in RAM for DBMS to operate on it!
  - Can't keep all the DBMS pages in main memory
- Buffer Manager: Efficiently uses main memory
  - Memory divided into **buffer frames**: slots for holding disk pages




❖ *Table of <frame#, pageid> pairs is maintained.*

# Example B+ Tree

- Search: Starting from root, examine index entries in non-leaf nodes, and traverse down the tree until a leaf node is reached
  - Non-leaf nodes can be searched using a binary or a linear search.
- Search for 5\*, 15\*, all data entries  $\geq 24^*$

Root




# Static Hashing

Number of primary pages  $N$  fixed, allocated sequentially

overflow pages may be needed when file grows

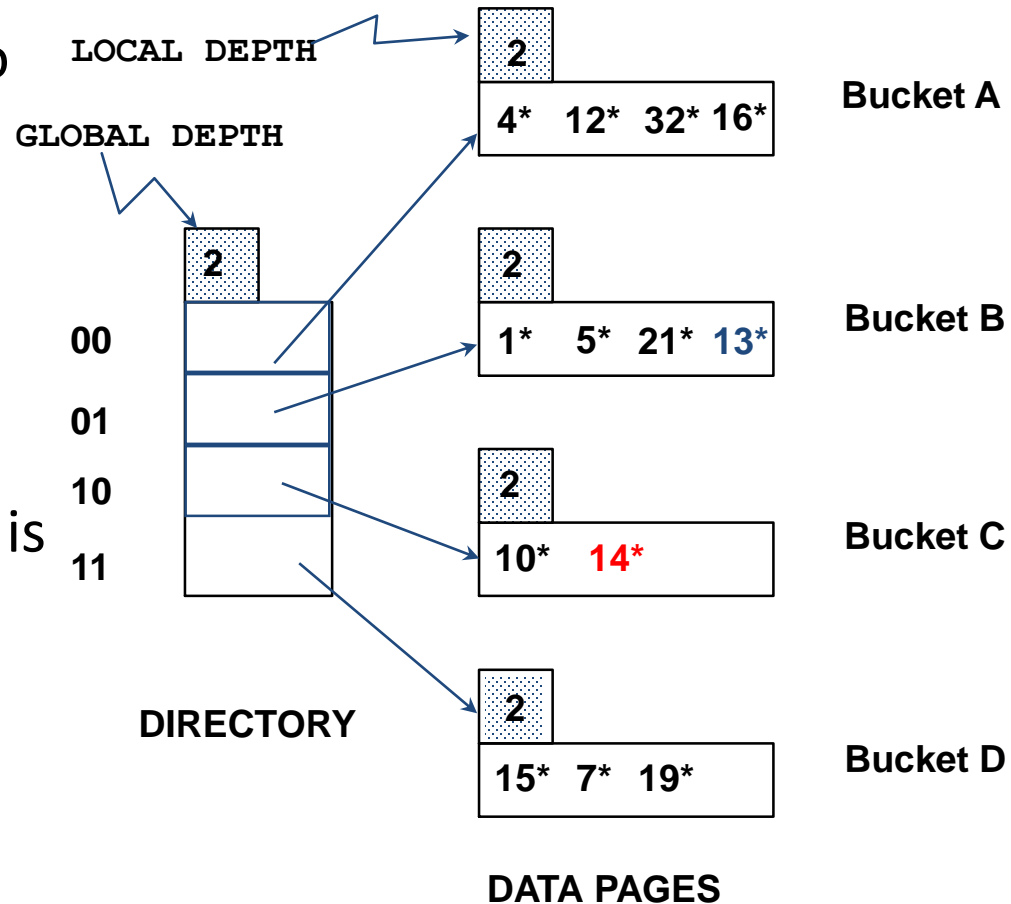
Buckets contain data entries

$h(k) \bmod N$  = bucket for data entry with key  $k$



# Extendible Hashing Example

- Directory is array of size 4
- Directory entry corresponds to last two bits of hash value
- If  $h(k) = 5 = \text{binary } 101$ , it is in bucket pointed to by 01
- Insertion into non-full buckets is trivial
- Insertion into full buckets requires split and **directory doubling**
- E.g., insert  $h(k)=20$



# Quad trees

- Nodes split at all dimensions at once
- Division fixed
- Cannot be balanced




Figure 5.17: A quad tree


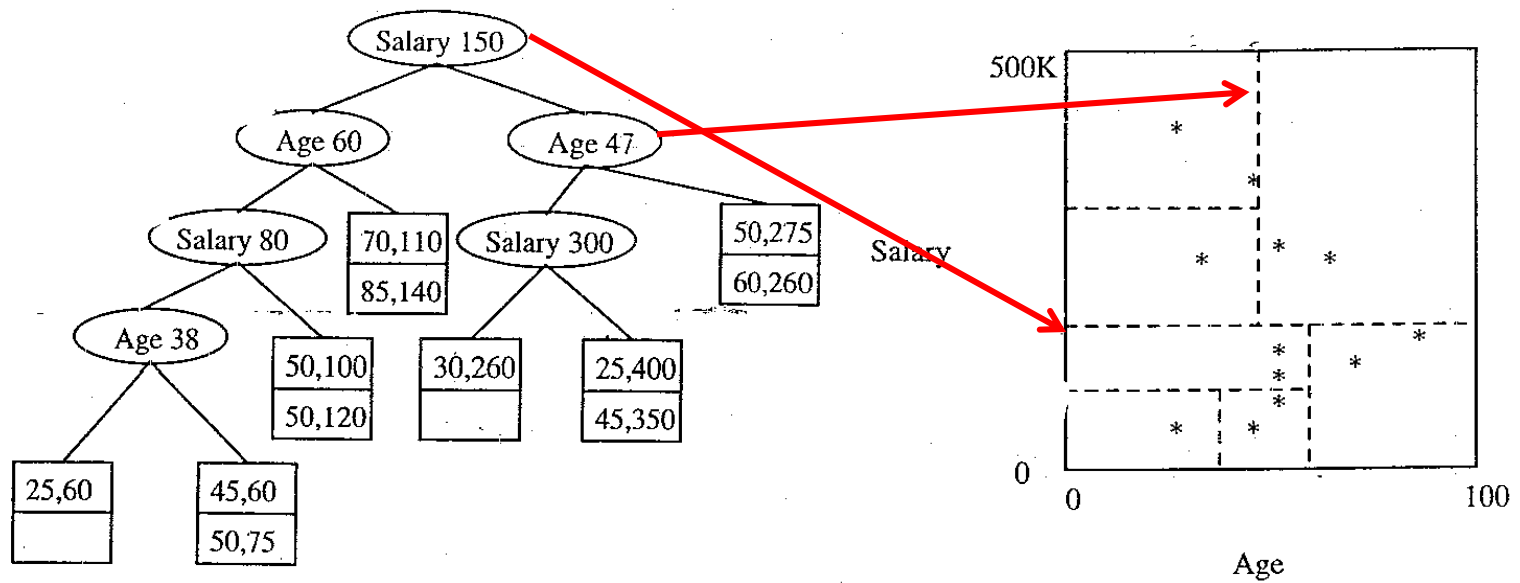


Figure 5.16: Data organized in a quad tree

# Kd-trees

- Each level of a  $k$ - $d$  tree partitions the space into two.
  - Choose one dimension for partitioning at the root level of the tree.
  - Choose another dimension for partitioning in nodes at the next level and so on, cycling through the dimensions.
- In each node, approximately half of the points stored in the sub-tree fall on one side and half on the other.
- Partitioning stops when a node has less than a given maximum number of points.

# KD-trees



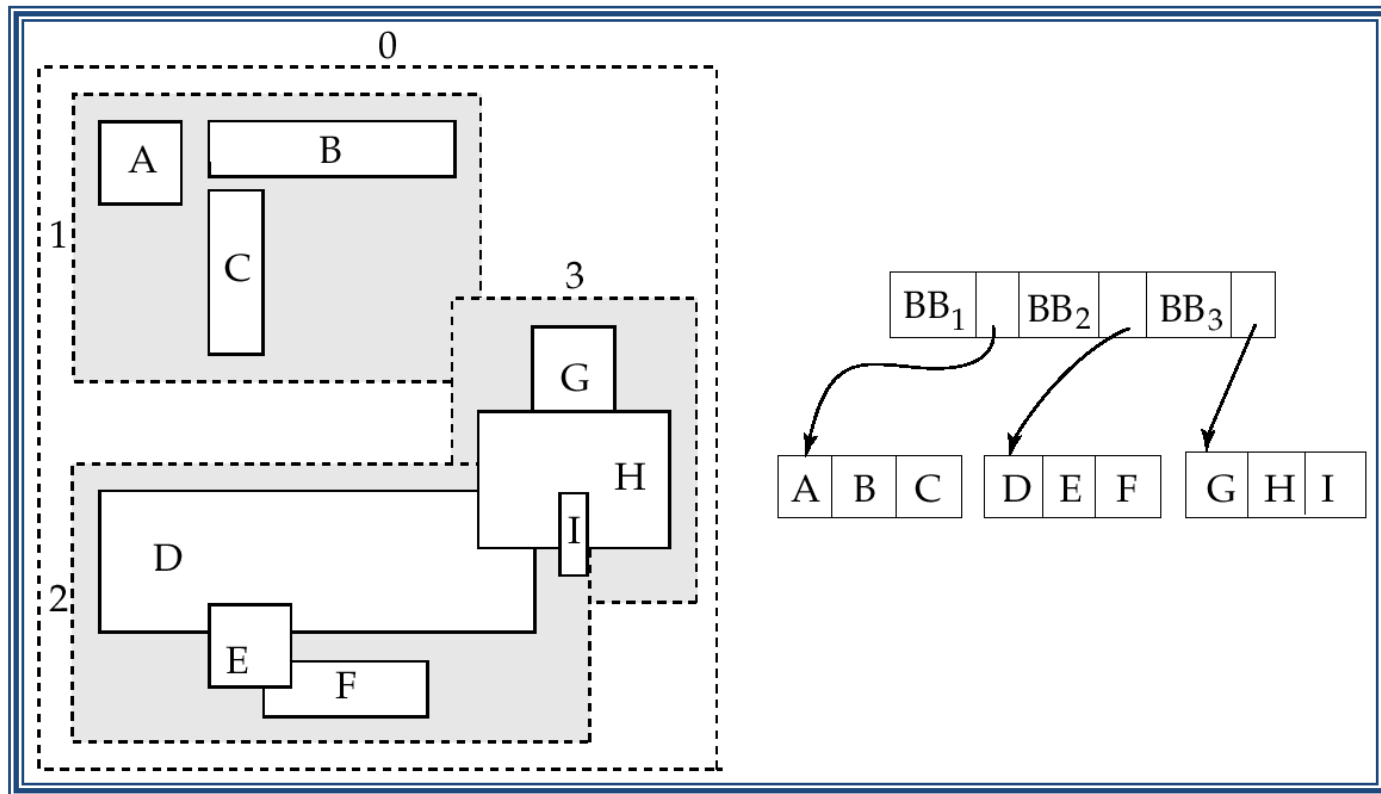
# R-Trees

- A rectangular **bounding box** is associated with each tree node.
  - Bounding box of a leaf node is a minimum sized rectangle that contains all the rectangles/polygons/regions associated with the leaf node.
  - The bounding box associated with a non-leaf node contains the bounding box associated with all its children.
  - Bounding box of a node serves as its key in its parent node (if any)
  - *Bounding boxes of children of a node are allowed to overlap*
- A polygon/region is stored only in one node, and the bounding box of the node must contain the polygon



# R-Trees: Example

A set of rectangles (solid line) and the bounding boxes (dashed line) of the nodes of an R-tree for the rectangles.



# FAQ: Why do we focus on IO cost?

## [1] CPU Cost

"Latency Numbers Every Programmer Should Know"

AM or check tuple equality in NLJ in RAM)

Typical algorithms in RAM (e.g., quicksort in  $n \log n$ )

It is hard for humans to get the picture until you translate it to "human numbers":

1 CPU cycle	1 s
Level 1 cache access	3 s
Level 2 cache access	9 s
Level 3 cache access	43 s
Main memory access	6 min
Solid-state disk I/O	2-6 days
Rotational disk I/O	1-12 months
Internet: SF to NYC	4 years
Internet: SF to UK	8 years
Internet: SF to Australia	19 years
OS virtualization reboot	423 years
SCSI command time-out	3000 years
Hardware virtualization reboot	4000 years
Physical system reboot	32 millenia

from HDD/SSDs  
(e read or write from HDD/SSD)

E.g., For ExternalSort

IO Cost ~=  
(sort N pages)

$$2N \left\lceil \log_B \frac{N}{2B} \right\rceil + 2N$$

data, focus on **IO cost** (i.e., it's the primary factor)  
< 10 GBs, just use RAM)

# Example NLJ vs. BNLJ: Steel Cage Match

Example:  $P(R) = 1000$ ,  $P(S) = 500$ ,  
 100 tuples/page  $\Rightarrow T(R) = 1000 \cdot 100$ ,  $T(S) = 500 \cdot 100$

	<b>B= 100 (+ 1 for output)</b>	<b>B= 20 (+ 1 for output)</b>
NLJ	$(1000 + 1000 \cdot 100 \cdot 500 + \text{OUT})$ $\Rightarrow \text{IO} = \sim 5,001,000 + \text{OUT}$	$(1000 + 1000 \cdot 100 \cdot 500 + \text{OUT})$ $\Rightarrow \text{IO} = \sim 5,001,000 + \text{OUT}$
BNLJ	$(1000 + 1000 \cdot 500 / 100)$ $\Rightarrow \text{IO} = \sim 6000 + \text{OUT}$	$(1000 + 1000 \cdot 500 / 20)$ $\Rightarrow \text{IO} = \sim 26,000 \text{ IOs} + \text{OUT}$

$$P(R) + T(R) \cdot P(S) + \text{OUT}$$

$$P(R) + P(R) \cdot P(S) / B + \text{OUT}$$

Small change in algorithm  $\Rightarrow$  Big speedup in JOINS ( **$\sim 1000\times$  faster**)  
 Also, notice if we swap R and S, we can save an extra **500 IOs** in BNLJ

# Example SMJ vs. BNLJ: Steel Cage Match

Consider  $P(R) = 1000$ ,  $P(S) = 500$

IO Cost  $\sim$  (sort N pages)

$$2N \left\lceil \log_B \frac{N}{2B} \right\rceil + 2N$$

	<b>B = 100</b>	<b>B = 20</b>
SMJ	(Sort R and S in NumPasses (np)=1 $2*1000*(np+1) + 2*500*(np+1) = 6000$ MergeJoin: $1000 + 500 = 1500$ IOs $\Rightarrow \text{IO} = 7500 \text{ IOs} + \text{OUT}$	(Sort R and S in NumPasses (np)=2 $2*1000*(np+1) + 2*500*(np+1) = 9000$ MergeJoin: $1000 + 500 = 1500$ IOs $\Rightarrow \text{IO} = 10,500 \text{ IOs} + \text{OUT}$
BNLJ	$(500 + 1000*500/100)$ $\Rightarrow \text{IO} = 5500 + \text{OUT}$	$(500 + 1000*500/20)$ $\Rightarrow \text{IO} = 25500 \text{ IOs} + \text{OUT}$

$$\sim \text{Sort}(P(R)) + \text{Sort}(P(S)) + P(R) + P(S) + \text{OUT}$$

$$P(R) + P(R)*P(S)/B + \text{OUT}$$

SMJ is  $\sim$  linear vs. BNLJ is quadratic...  
 Redo the same with 10x? SMJ much faster.

# Join Algorithms: Summary

For  $R \bowtie S$  on column  $A$

- NLJ: An example of a *non*-IO aware join algorithm
- BNLJ: Big gains just by being IO aware & reading in chunks of pages!

Quadratic in  $P(R), P(S)$   
I.e.  $O(P(R)*P(S))$

- 
- SMJ: Sort  $R$  and  $S$ , then scan over to join!
  - HPJ: Partition  $R$  and  $S$  into buckets using a hash function, then join the (much smaller) matching buckets

Given sufficient buffer space, linear in  $P(R), P(S)$   
I.e.  $\sim O(P(R)+P(S))$

By only supporting equijoins & taking advantage of this structure!